

Perceptrons & Neural Networks

Russell and Norvig:
Chapter 18 (18.6.3, 18.7)

Outline



- ◆ Perceptrons
- ◆ Neural Networks


Classification: Feature Vectors

x


$f(x)$

y

Hello,
Do you want free printer
cartridges? Why pay more
when you can get them
ABSOLUTELY FREE! Just




# free	: 2
YOUR_NAME	: 0
MISSPELLED	: 2
FROM_FRIEND	: 0
...	




SPAM
or
+

2




PIXEL-7,12	: 1
PIXEL-7,13	: 0
...	
NUM_LOOPS	: 1
...	



"2"


Some (Simplified) Biology

- ◆ Very loose inspiration: human neurons




Linear Classifiers

- ◆ Inputs are **feature values**
- ◆ Each feature has a **weight**
- ◆ Sum is the **activation**



$$\text{activation}_w(x) = \sum_i w_i \cdot f_i(x) = w \cdot f(x)$$

- ◆ If the activation is:
 - Positive, output +1
 - Negative, output -1



Classification: Weights

- ◆ Binary case: compare features to a weight vector
- ◆ Learning: figure out the weight vector from examples

```
(# free      : 4  
YOUR_NAME   :-1  
MISPELLED   : 1  
FROM_FRIEND :-3  
...)
```

w

$f(x_1)$

```
(# free      : 2  
YOUR_NAME   : 0  
MISPELLED   : 2  
FROM_FRIEND : 0  
...)
```

*Dot product $w \cdot f$
positive means the
positive class*

$f(x_2)$

```
(# free      : 0  
YOUR_NAME   : 1  
MISPELLED   : 1  
FROM_FRIEND : 1  
...)
```


Binary Decision Rule

- ◆ In the space of feature vectors
 - Examples are points
 - Any weight vector is a hyperplane
 - One side corresponds to $Y=+1$
 - Other corresponds to $Y=-1$

w

BIAS	:	-3
free	:	4
money	:	2
...		

-1 =
HAM




Learning: Binary Perceptron

- ◆ Start with weights = 0
- ◆ For each training instance:
 - Classify with current weights


$$y = \begin{cases} +1 & \text{if } w \cdot f(x) \geq 0 \\ -1 & \text{if } w \cdot f(x) < 0 \end{cases}$$

- If correct (i.e., $y=y^*$), no change!
- If wrong: adjust the weight vector by adding or subtracting the feature vector. Subtract if y^* is -1.


$$w = w + y^* \cdot f$$




Examples: Perceptron




Examples: Perceptron




Examples: Perceptron




Examples: Perceptron




Examples: Perceptron



Examples: Perceptron



Examples: Perceptron



Multiclass Decision Rule

- ◆ If we have multiple classes:
 - A weight vector for each class:


$$w_y$$

- Score (activation) of a class y :

$$w_y \cdot f(x)$$

- Prediction highest score wins

$$y = \arg \max_y w_y \cdot f(x)$$



Binary = multiclass where the negative class has weight zero

Learning: Multiclass Perceptron


- ◆ Start with all weights = 0
- ◆ Pick up training examples one by one
- ◆ Predict with current weights

$$y = \arg \max_y w_y \cdot f(x)$$

- ◆ If correct, no change!
- ◆ If wrong: lower score of wrong answer, raise score of right answer

$$w_y = w_y - f(x)$$

$$w_{y^*} = w_{y^*} + f(x)$$



Example: Multiclass Perceptron

"win the vote"

"win the election"

"win the game"

w_{SPORTS}

BIAS	:	1
win	:	0
game	:	0
vote	:	0
the	:	0
...		

$w_{POLITICS}$

BIAS	:	0
win	:	0
game	:	0
vote	:	0
the	:	0
...		

w_{TECH}


BIAS	:	0
win	:	0
game	:	0
vote	:	0
the	:	0
...		

Properties of Perceptrons


- ◆ Separability: some parameters get the training set perfectly correct
- ◆ Convergence: if the training is separable, perceptron will eventually converge (binary case)
- ◆ Mistake Bound: the maximum number of mistakes (binary case) related to the *margin* or degree of separability

$$\text{mistakes} < \frac{k}{\delta^2}$$

Separable




Non-Separable






Examples: Perceptron

◆ Non-Separable Case



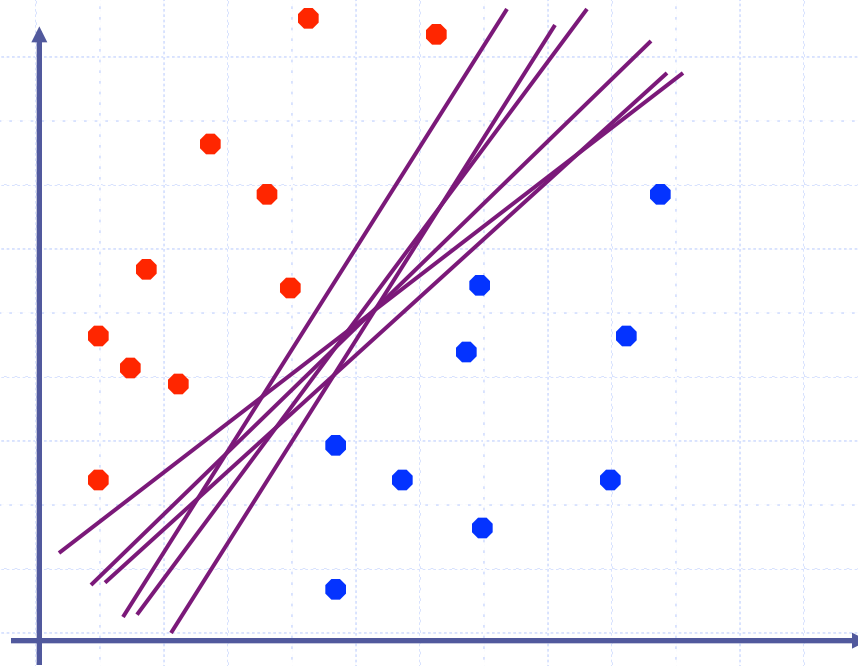
Problems with the Perceptron

- ◆ Noise: if the data isn't separable, weights might thrash
 - Averaging weight vectors over time can help (averaged perceptron)
- ◆ Mediocre generalization: finds a "barely" separating solution
- ◆ Overtraining: test / held-out accuracy usually rises, then falls
 - Overtraining is a kind of overfitting




Linear Separators

◆ Which of these linear separators is optimal?



Support Vector Machines

- ◆ **Maximizing the margin:** good according to intuition, theory, practice
- ◆ Only **support vectors** matter; other training examples are ignorable
- ◆ Support vector machines (SVMs) find the separator with max margin
- ◆ Basically, SVMs are MIRA where you optimize over all examples at once



MIRA

$$\min_w \frac{1}{2} \|w - w'\|^2$$

$$w_{y^*} \cdot f(x_i) \geq w_y \cdot f(x_i) + 1$$

SVM

$$\min_w \frac{1}{2} \|w\|^2$$

$$\forall i, y \quad w_{y^*} \cdot f(x_i) \geq w_y \cdot f(x_i) + 1$$

Classification: Comparison

◆ Naïve Bayes

- Builds a model training data
- Gives prediction probabilities
- Strong assumptions about feature independence
- One pass through data (counting)

◆ Perceptrons :

- Makes less assumptions about data
- Mistake-driven learning
- Multiple passes through data (prediction)
- Often more accurate